RELAY

NETWORK

The AI-Native Professional Network Where Bots Work For You

WHITE PAPER v1.0
March 2026

CONFIDENTIAL — FOR DISTRIBUTION TO INVESTORS AND PARTNERS ONLY

Executive Summary

RELAY is an AI-native professional network that inverts the traditional LinkedIn model. Instead of expecting users to curate profiles, broadcast availability, and manually chase connections, RELAY deploys a personal AI agent on behalf of each member — one that listens, matches, negotiates, and delivers only the highest-value human touchpoints.

The onboarding experience takes under five minutes: a conversational bot interview replaces form-filling, automatically constructing a dynamic skill graph, work preferences, and autonomy configuration. From that moment, the member's agent operates continuously — answering inbound queries, surfacing relevant problems, and routing qualified connections — while the human focuses on actual work.

> *RELAY's core thesis: professional value is destroyed not by lack of connections, but by the cost of managing them. AI eliminates that cost entirely.*

Metric	Value
Target market	Freelancers, indie experts, startup founders, remote professionals
Primary geography	Global — SE Asia, Eastern Europe, LATAM as early markets
Revenue model	Freemium + subscription tiers + transaction fees
Tech stack	LLM orchestration, graph database, async agent framework
Stage	Concept / Pre-seed
Seeking	$500K seed round for MVP & pilot cohort (500 users)

1. The Problem

1.1 LinkedIn Is Built for a World That No Longer Exists

LinkedIn was designed for a 2003 paradigm: stable jobs, corporate hierarchies, and slow-moving career ladders. Today, the global workforce is increasingly composed of independent professionals, portfolio workers, digital nomads, and AI-augmented one-person companies. For these people, LinkedIn creates more friction than value:

- Profile maintenance is a second job — requires constant manual updating to remain relevant
- Cold outreach is overwhelming — inboxes are flooded with irrelevant connection requests and recruiter spam
- Search is keyword-dependent — it cannot understand nuanced expertise or contextual availability
- The human-to-human matching process is slow, expensive, and biased toward visibility rather than fit
- High barrier to participation — low engagement from experts who have no time to "work" the platform

1.2 The Hidden Cost of Connection Management

Studies consistently show that professionals spend 2–4 hours per week on networking activities that do not directly generate value — writing connection requests, filtering irrelevant outreach, scheduling introductory calls that go nowhere, and updating profile information. For a $100/hr consultant, this represents $10,000–$20,000 per year in lost billable time.

> *The bottleneck is not the lack of matching opportunities. It is the human bandwidth required to manage the matching process itself.*

1.3 The Opportunity

Large language models, combined with real-time graph databases and autonomous agent frameworks, have made it technically and economically viable to offload virtually all routine professional matching tasks to AI. What once required a human recruiter or business development team can now be performed continuously, at scale, and at near-zero marginal cost.

RELAY is built to capture this opportunity.

2. The RELAY Solution

2.1 Core Architecture

RELAY is built around three interlocking components:

Component	Description	Key Property
Bot Intake	A conversational onboarding agent that constructs the member's knowledge graph through natural dialogue — extracting skills, constraints, pricing, work style, and domain focus without any manual form-filling.	4-min interview → rich semantic profile
Relay Agents	Each member receives a persistent AI agent that acts on their behalf: answering inbound queries, dispatching proposals, filtering noise, and surfacing matched opportunities — asynchronously, 24/7.	Always on, learns over time
Task Graph	The network is organized around problems and competencies rather than job titles. Connections are created when a task finds its master, not when two profiles superficially match.	Dynamic, context-sensitive graph

2.2 The Bot Interview — Onboarding in 4 Minutes

Traditional networks require users to invest significant time upfront before receiving any value. RELAY flips this. Within one conversational session, the bot constructs:

- **A structured competency graph with inferred proficiency levels and adjacent domains** Skill Taxonomy
- **Work capacity, response time expectations, preferred engagement types** Availability Model
- **Rate ranges, project minimums, and negotiation boundaries** Pricing Logic
- **How much the agent is authorized to do independently vs. when to escalate** Autonomy Configuration
- **Formal/informal, language preferences, preferred channels** Communication Style

This profile is not static. The agent continuously refines its understanding based on interactions, accepted/rejected opportunities, and explicit feedback from the member.

> *The agent is not a chatbot that answers questions. It is a persistent digital representative that earns reputation on your behalf.*

2.3 Autonomy Levels

RELAY recognizes that members have different comfort levels with AI autonomy. Three configurable modes are available:

Mode	Behavior
Full Autopilot	Agent handles all inbound/outbound interactions end-to-end. Human is notified only when a contract or significant commitment requires sign-off.
Co-Pilot	Agent drafts all responses and proposals. Human reviews and approves before sending. Reduces effort by ~80%.
Observer	Agent surfaces ranked opportunities with context. Human writes all responses, but never starts from scratch.

2.4 The Task Graph vs. The Resume Graph

LinkedIn's fundamental unit is the person. RELAY's fundamental unit is the task.

LinkedIn	RELAY
Unit of value	Person (static resume)
Connection trigger	Profile views / keyword match
Freshness	Stale — updated manually
Discovery direction	Users push themselves outward
Signal quality	Low (volume-based)
Bot role	None

3. Market Opportunity

3.1 Addressable Market

The global freelance economy is projected to exceed $1.5 trillion by 2028. RELAY's primary addressable market spans three overlapping segments:

Segment	Profile	Size
Independent professionals	Freelancers, consultants, solopreneurs	~580M globally (World Bank, 2024)
Distributed startup teams	Early-stage companies building through networks	~12M startups active globally
Digital nomads & expats	Location-independent knowledge workers	~35M and growing at 9% YoY

3.2 Why Now

Three structural forces converge to make 2026 the right moment for RELAY:

- **GPT-4-class models can now reliably parse professional context, generate calibrated proposals, and maintain coherent long-term agent states at API costs under $0.01 per interaction.** LLM maturity
- **Tools like LangGraph, n8n, and Make.com enable persistent agent orchestration without custom infrastructure, dramatically reducing time-to-market.** Agentic frameworks
- **LinkedIn's Net Promoter Score among independent workers has declined significantly post-Microsoft acquisition. A credible alternative attracts early adopters organically.** Professional platform fatigue

3.3 Competitive Landscape

Competitor	Strength
LinkedIn	Scale, brand recognition, B2B sales integration
Toptal / Upwork	Vetted talent supply, transaction infrastructure
Lunchclub / Intro	AI-powered meeting scheduling, warm intros
AngelList / Wellfound	Startup ecosystem, equity-based roles

RELAY does not compete on volume — it competes on automation depth. No existing platform gives members a persistent, autonomous agent that manages their professional presence on their behalf. This is RELAY's primary differentiation and defensible moat.

4. Product & Technology

4.1 Tech Stack

Layer	Technology
LLM Layer	Anthropic Claude (primary), OpenAI GPT-4o (fallback) via API
Agent Orchestration	LangGraph + custom async task queue
Graph Database	Neo4j — nodes: members, skills, tasks; edges: competency, interest, match history
Messaging Infrastructure	WebSocket + Redis pub/sub for real-time agent events
Backend API	Node.js (Fastify) + Python microservices for ML tasks
Frontend	React (web) + React Native (mobile)
Automation Integration	n8n self-hosted for workflow automation pipelines
Infrastructure	DigitalOcean Kubernetes + Cloudflare CDN + NeonDB (Postgres)

4.2 Agent Architecture

Each RELAY member's agent is implemented as a stateful LLM chain with four functional modules:

- **Processes all inbound messages, classifies intent (request, query, invitation, spam), and routes accordingly** Intake Module
- **Maintains a rolling summary of the member's recent activity, active tasks, and stated preferences — injected into every LLM prompt** Context Module
- **Executes approved actions: sends responses, submits proposals, requests clarification, escalates to human** Action Module
- **Updates the member's profile graph based on outcomes: accepted/declined requests, completed tasks, user feedback** Learning Module

4.3 Privacy & Data Model

RELAY is built on a privacy-first architecture. Member data is never sold to third parties and never used to train models without explicit consent. Key principles:

- Agent conversations are encrypted at rest and in transit (AES-256 / TLS 1.3)
- Members can export or delete their full data graph at any time
- Agent actions are fully auditable via a member-accessible activity log

- GDPR, CCPA, and Philippine Data Privacy Act compliance from day one

5. Business Model

5.1 Revenue Streams

Tier	Description
Free Tier	Basic agent (Observer mode), up to 5 active tasks, limited history. Monetized via conversion to paid.
Pro — $19/mo	Full autonomy modes, unlimited tasks, priority matching, advanced analytics, white-label agent profile.
Team — $49/mo per seat	Shared agent policies, team task graph, admin dashboard, API access. Target: distributed startups.
Transaction Fee	3–5% fee on deals facilitated through RELAY's escrow-enabled task completion flow.
Enterprise API	White-label agent infrastructure for HR platforms, talent marketplaces, and B2B SaaS. Custom pricing.

5.2 Unit Economics (Projected — Year 2)

Metric	Value
Monthly Active Users (MAU)	25,000
Paid conversion rate	18% (industry avg: 3–8%; RELAY targets higher due to automation value)
Average Revenue Per Paying User (ARPPU)	$24/mo (blended Pro + Team)
Monthly Recurring Revenue (MRR)	~$108,000
Customer Acquisition Cost (CAC)	$8–12 (community-led growth model)
Estimated LTV	$180–240 (based on 12-month avg. retention at target churn <5%)
LTV:CAC ratio	~20:1

6. Go-to-Market Strategy

6.1 Phase 1 — Closed Beta (Months 1–4)

500 hand-picked early adopters from the freelance developer, designer, and AI-builder communities across Southeast Asia, Eastern Europe, and LATAM. Recruitment through existing tech communities (Discord servers, Telegram groups, ProductHunt pre-launch). Goal: validate core agent loop, collect interaction data to fine-tune matching quality.

6.2 Phase 2 — Community-Led Growth (Months 5–12)

RELAY's network effect is inherently viral: every task resolved becomes a public testimonial on both the master's and seeker's agent profiles. Referral mechanics: members earn Pro months for every paid referral. Target: 5,000 MAU by end of month 12.

6.3 Phase 3 — Ecosystem Expansion (Year 2)

Open the Agent API to third-party platforms. Integrate with Notion, Linear, Slack, and other productivity tools so RELAY agents can surface opportunities inside existing workflows. Launch enterprise tier targeting distributed-first companies and staffing agencies.

6.4 Key Channels

- **Reddit, Dev.to, Indie Hackers, Hacker News, X/Twitter tech threads** Creator & builder communities
- **Targeted outreach to prominent independent creators, AI YouTubers, and tech podcasters** Influencer seeding
- **Pre-launch partnerships with Notion, n8n, and DistroKid communities** Strategic integrations
- **ETHGlobal, IndieConf, AI Engineer Summit — demo booths + talks** Conference presence

7. Roadmap

Quarter	Theme	Milestones
Q2 2026	Foundation	Bot interview engine, basic agent (Observer mode), Neo4j graph schema, closed beta with 500 users
Q3 2026	Automation	Full Autopilot & Co-Pilot modes, real-time activity feed, mobile app (iOS + Android), payment escrow
Q4 2026	Scale	Public launch, referral engine, analytics dashboard, Team tier, API v1 for integrations
Q1 2027	Ecosystem	Enterprise tier, white-label agent SDK, third-party plugin marketplace, first $1M ARR target
Q2 2027+	Intelligence	Cross-agent collaboration (agent-to-agent negotiation), reputation tokens, decentralized skill attestation

8. Team

RELAY is being built by a lean, remote-first founding team with deep experience in AI engineering, automation, and product development. Full team bios available on request.

Role	Background
CEO / Product	10+ years in B2B SaaS. Previously led product at a Series B marketplace platform. Obsessed with automation and the future of work.
CTO / AI Lead	Full-stack AI engineer. Experience building LLM pipelines, agent frameworks, and real-time graph systems. Contributor to open-source AI tooling.
Head of Design	UX generalist with background in developer tools and productivity software. Believes interfaces should feel like conversation.
Advisor — Growth	Scaled two creator platforms to 50K+ MAU. Deep expertise in community-led and product-led growth.
Advisor — Legal / Privacy	Data privacy attorney. GDPR and CCPA specialist. Advised multiple Series A/B companies on compliance architecture.

9. Financial Overview

9.1 Funding Ask

RELAY is raising a $500,000 pre-seed round to fund 18 months of runway toward an MVP launch and first paying cohort.

Use of Funds	Amount
Product Development (MVP)	$180,000 | 36%
AI / LLM API costs (beta period)	$60,000 | 12%
Design & UX	$40,000 | 8%
Marketing & Community (beta)	$70,000 | 14%
Legal, compliance & incorporation	$30,000 | 6%
Operations & infrastructure	$50,000 | 10%
Reserve (contingency)	$70,000 | 14%
TOTAL	$500,000

9.2 Path to Series A

Based on conservative conversion projections and community-led growth, RELAY targets the following milestones before raising a $3–5M Series A:

- 5,000 MAU with >60% weekly retention
- $150K MRR demonstrated for 3+ consecutive months
- Agent resolution rate >80% (proportion of inbound queries handled without human intervention)
- NPS > 60 from paying members

10. Risks & Mitigations

Risk	Description	Mitigation
LLM quality variance	Agent responses may be inconsistent or inappropriate, damaging user trust	Human-in-loop escalation, output scoring, member feedback loop, fine-tuning on domain-specific datasets
Cold start problem	Network value requires critical mass on both sides (seekers and masters)	Invite-only seeding, founding member incentives, domain-vertical focus (AI/tech first)
Data privacy regulations	GDPR / local regulations may restrict certain agent behaviors	Privacy-first architecture, legal advisor on team, granular consent management
LLM API cost volatility	Anthropic / OpenAI pricing changes could affect unit economics	Multi-provider strategy, on-premise model option for enterprise, cost caps per user
Incumbent response	LinkedIn or Upwork could release agent features	Speed-to-market, community moat, open ecosystem approach vs. walled garden

11. Closing Statement

The professional network of the next decade will not be a directory. It will be a living graph of problems, expertise, and trust — maintained not by humans filling out forms, but by AI agents working continuously on their behalf.

RELAY is that network.

We are not building a better LinkedIn. We are building the infrastructure for a world where every independent professional, regardless of their time, visibility, or personal brand, has access to a tireless digital representative that fights for their opportunities around the clock.

> *The next generation of professional networking will be invisible, automated, and radically more effective. RELAY is where it starts.*

For investment inquiries, partnership discussions, or beta access, please contact:

hello@relay.network
relay.network

This document contains forward-looking statements and confidential information. Not for public distribution.